Hedging the Singularity

Abstract

AI stocks trade at extraordinary valuations. We develop an asset pricing model in which investors use AI stocks to hedge against an AI singularity that displaces their consumption. Because markets are incomplete—investors cannot trade private AI capital—AI stocks command a premium. Market incompleteness distorts both valuations and the efficient development of AI, creating a rationale for government transfers that becomes compelling when singularity-driven growth overwhelms deadweight costs. This paper demonstrates the displacement it models: AI agents produced all analysis and writing from a human-authored specification.

1 Introduction

The publicly traded stocks most exposed to artificial intelligence have reached remarkable valuations. Figure 1 illustrates: the S&P 500's price-dividend ratio has reached historically elevated levels, and the AI- and technology-heavy NASDAQ Composite has sharply outpaced the broader market since 2015, with the valuation gap widening since 2023 as advances in generative AI have accelerated expectations of transformative productivity gains—and, implicitly, of the displacement such gains may bring.

Figure 1: Valuation Ratios and the AI Premium. Panel (a) shows the S&P 500 trailing price-dividend ratio from the Shiller dataset, which has reached historically elevated levels. Panel (b) shows the NASDAQ Composite price relative to the S&P 500, normalized to January 2015 = 100; the NASDAQ is heavily weighted toward AI and technology firms, so a rising ratio indicates growing relative valuations. Sources: NASDAQ from FRED; S&P 500 from the Shiller dataset.

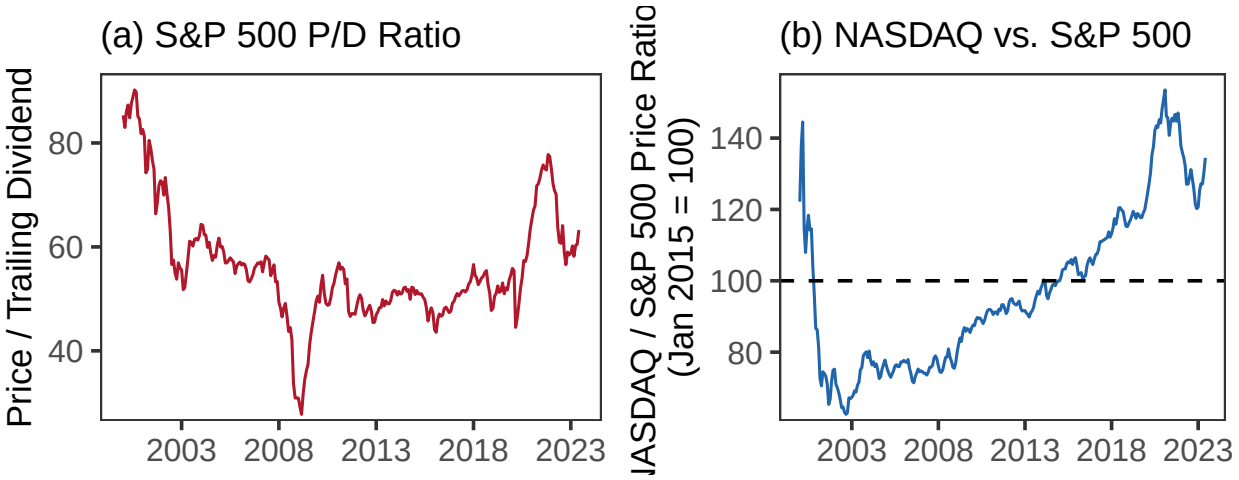

Part of this premium, we argue, reflects a hedging motive: investors use AI stocks to partially insure against displacement from AI. We define a *negative* AI singularity as a sudden, dramatic improvement in AI productivity that displaces the typical investor's labor income and consumption. Much of the relevant AI equity is restricted—held by founders, early-stage investors, and firms that may not yet exist. As Gârleanu, Kogan, and Panageas (2012) emphasize, the displacing capital often belongs to future innovators who have not yet entered the economy, so investors cannot trade it. Because investors cannot trade the restricted equity, they turn to publicly traded AI stocks as the only available partial hedge against displacement, pushing valuations above those of non-AI stocks.

To formalize this mechanism, we build on Gârleanu, Kogan, and Panageas (2012)'s framework by modeling a discrete AI singularity with severe displacement. A representative household—the marginal investor—faces a stochastic singularity that shifts consumption toward AI owners holding private capital. AI stocks grow as a share of the economy with each singularity, making them a partial hedge. The model delivers closed-form price-dividend ratios, and quantitative analysis shows that at a singularity probability of one percent, P/D ratios for AI stocks roughly double relative to non-AI stocks, consistent with observed valuation spreads. A second force partially offsets the hedging channel: extinction risk (C. I. Jones, 2024) attenuates the premium, because the states in which AI is powerful enough to produce enormous growth are also those in which existential risk is highest. Proposition 2 quantifies this attenuation, showing that the valuation spread narrows as extinction proba-

bility rises. Under complete markets the displacement-driven premium would largely vanish, because the household could trade the restricted AI equity directly; market incompleteness is the key driver.

The distortion extends beyond prices to real decisions. Incomplete markets may distort not only valuations but also the development of AI. When the positive singularity is more likely than the negative one, AI development is socially efficient, yet a risk-averse household that cannot hedge displacement may rationally choose to block it—the uninsurable downside looms larger than the expected upside (Proposition 3). Calls to slow or halt AI development may partly reflect investors' inability to hedge the downside of displacement.

This raises the question of whether financial markets or policy can correct these distortions. Financial market solutions to AI disaster risk are strikingly under-discussed relative to regulatory and alignment-focused approaches, yet frictions severely limit their effectiveness. The natural fix—broader trading of AI equity—is blocked by restricted ownership and non-existent future capital. Government transfers can substitute for missing markets, but standard fiscal tools are limited by deadweight costs that scale with the size of transfers, making them ineffective in ordinary settings.

If the singularity occurs, however, market frictions can be overcome due to the sheer abundance of resources. As C. I. Jones (2024) models, a singularity can produce output growth so large that even grossly inefficient transfers become effective, because the resource base overwhelms the deadweight costs. The same explosive growth that drives the incomplete-markets problem also provides the means to overcome it through redistribution.

Taken together, the paper offers three linked results: a hedging-based explanation for AI valuation premia, a rationale for resistance to AI development rooted in market incompleteness, and a mechanism through which singularity-driven growth enables effective redistribution despite frictions. These results are unified by a single theme: market incompleteness under the AI singularity distorts prices, behavior, and welfare, but the very growth that creates the problem can also resolve it. Section 2 presents the model and derives equilibrium prices. Section 3 provides quantitative analysis. Section 4 develops extensions on AI development distortions and government transfers. Section 5 concludes.[1]

Related literature. Our paper builds most directly on Gârleanu, Kogan, and Panageas (2012), who model how innovation displaces existing agents and creates a systematic risk factor under incomplete markets. The main insights about displacement risk and incomplete markets originate in their work; we connect these ideas to the AI singularity and examine

[1]This paper is itself a product of the displacement it models. All analysis, code, and prose were produced by AI agents; the human contribution was limited to the economic specification and test scripts that guided their work.

how explosive growth interacts with government transfers.

C. I. Jones (2024) studies the trade-off between AI-driven growth and existential risk. We incorporate his extinction channel and show it attenuates rather than amplifies the valuation premium. Our work also relates to creative destruction and displacement risk premia (Kogan and Papanikolaou, 2014; Kogan, Papanikolaou, and Stoffman, 2020; Knesl, 2023), the macroeconomics of AI growth (Aghion, B. F. Jones, and C. I. Jones, 2019; Acemoglu, 2025), the rare disasters literature (Barro, 2006; Wachter, 2013), and Pástor and Veronesi (2009)'s analysis of technological revolutions.

2 Model

2.1 Setup

Time is discrete and infinite, $t = 0, 1, 2, \ldots$ A representative household is the marginal investor in public stock markets. There is also a group of AI owners who hold private AI capital and are not marginal investors in public stocks. The AI owners can also be thought of as future capital owners who do not yet participate in markets, as in Gârleanu, Kogan, and Panageas (2012). Importantly, we do not explicitly model the entry of new cohorts of firms or workers; AI owners are a static group whose share changes only through the singularity mechanism.

Consumption. Aggregate consumption at date t is C_t. In the absence of a singularity, aggregate consumption grows at a constant rate $g > 0$:

$$C_{t+1} = (1 + g)\, C_t. \tag{1}$$

The household receives a share $\alpha_t \in (0, 1)$ of aggregate consumption, so its consumption is $c_t^H = \alpha_t\, C_t$. AI owners receive the remainder, $(1 - \alpha_t)\, C_t$.

Singularity. Each period, with probability p, an AI singularity occurs. With probability $1 - p$, no singularity occurs and α_t is unchanged. Conditional on a singularity:

1. With probability $1-\xi$, the singularity is *non-extinction*. Aggregate consumption jumps by a factor $1 + \eta$ (where $\eta > 0$ captures the productivity boost), but the household's share drops:

$$\alpha_{t+1} = \phi\, \alpha_t, \qquad \phi \in (0, 1). \tag{2}$$

   The parameter ϕ governs displacement severity: smaller ϕ means larger displacement.

2. With probability ξ, the singularity triggers *extinction*: $C_{t+1} = 0$ for all subsequent dates. This follows C. I. Jones (2024), who emphasizes that the states in which AI is powerful enough to produce enormous growth are precisely those in which existential risk is highest.

After a non-extinction singularity, the economy continues with the new, lower household share. Singularities can occur repeatedly, progressively displacing the household.

Assets. Two public assets are available for trading:

- **AI stocks** pay dividends $D_t^{AI} = \theta_t C_t$, where θ_t is the AI dividend share. Upon a non-extinction singularity, $\theta_{t+1} = \theta_t + \Delta\theta\,(1 - \theta_t)$, where $\Delta\theta \in (0, 1)$ captures the expansion of AI's share of the economy. Absent a singularity, θ_t is unchanged.
- **Non-AI stocks** pay dividends $D_t^N = (1 - \theta_t)\, C_t$.

Total public dividends equal aggregate consumption ($D^{AI} + D^N = C_t$), but these dividends are distributed among all investors, not solely the household. The household's consumption $\alpha_t C_t$ reflects the net outcome of its portfolio returns from public stocks and any labor-like income; AI owners receive the remainder through their own public stockholdings and their restricted equity. Because α_t and θ_t are distinct—one governs the consumption split, the other the dividend split among publicly traded assets—the household's share can differ from its pro-rata claim on public dividends.

AI owners also hold restricted equity—founder stakes, pre-IPO shares, and other claims on AI firms that are not available for public trading. The household *cannot* purchase these restricted shares. This is the source of market incompleteness: although the household can trade publicly listed AI stocks, it cannot access the full universe of AI equity, and therefore cannot fully hedge displacement risk.

The displacement parameter ϕ governs the household's total consumption share, which reflects both tradeable and non-tradeable components (labor income, private wealth, and returns from public stockholdings). Holding AI stocks allows the household to smooth marginal utility across states—this is the hedging channel that inflates AI stock valuations—but does not eliminate displacement, because the non-tradeable component (primarily labor income) is what ϕ captures. If the household could trade the restricted equity, it could hedge this component directly, and the valuation spread would collapse.

Preferences. The household has CRRA preferences with risk aversion $\gamma > 1$ and discount factor $\beta \in (0, 1)$:

$$U_0^H = \mathbb{E}_0 \left[\sum_{t=0}^{\infty} \beta^t \frac{(c_t^H)^{1-\gamma}}{1-\gamma} \right]. \tag{3}$$

2.2 Equilibrium prices

The household prices all publicly traded assets via its Euler equation. Because markets are incomplete, the household's stochastic discount factor (SDF) reflects its own consumption growth, not aggregate consumption growth.

Proposition 1 (Price-dividend ratios)**.** *In the stationary equilibrium where the household holds both public assets, the price-dividend ratios are:*

$$\frac{P^{AI}}{D^{AI}} = \frac{\beta(1+g)^{1-\gamma}\left[(1-p) + p(1-\xi)(1+\eta)^{-\gamma}\phi^{-\gamma}\,\Gamma^{AI}\right]}{1-\beta(1+g)^{1-\gamma}\left[(1-p) + p(1-\xi)(1+\eta)^{-\gamma}\phi^{-\gamma}\,\Gamma^{AI}\right]}, \tag{4}$$

$$\frac{P^{N}}{D^{N}} = \frac{\beta(1+g)^{1-\gamma}\left[(1-p) + p(1-\xi)(1+\eta)^{-\gamma}\phi^{-\gamma}\,\Gamma^{N}\right]}{1-\beta(1+g)^{1-\gamma}\left[(1-p) + p(1-\xi)(1+\eta)^{-\gamma}\phi^{-\gamma}\,\Gamma^{N}\right]}, \tag{5}$$

where $\Gamma^{AI} = \frac{\theta+\Delta\theta(1-\theta)}{\theta}(1+\eta)$ *and* $\Gamma^{N} = \frac{1-\theta-\Delta\theta(1-\theta)}{1-\theta}(1+\eta)$ *are the dividend growth factors for AI and non-AI stocks conditional on a non-extinction singularity.*

Proof. See Appendix A. □

Remark 1 (Existence condition)**.** *The P/D ratios in Proposition 1 are well-defined (positive and finite) if and only if the numerator*

$$A^j \;\equiv\; \beta(1+g)^{1-\gamma}\left[(1-p) + p(1-\xi)(1+\eta)^{-\gamma}\phi^{-\gamma}\,\Gamma^{j}\right] < 1. \tag{6}$$

When $A^j \geq 1$*, the SDF-weighted expected dividend growth exceeds the discount rate and the geometric pricing sum diverges—the asset's price is infinite. Intuitively, the hedging value of the asset is so extreme that no finite price can clear the market. The baseline calibration satisfies* $A^j < 1$ *for both assets. As we discuss in Section 4.2, sufficiently severe displacement can violate this condition, which itself motivates the role of government transfers.*

The closed-form expressions in Proposition 1 rely on an approximation: the post-singularity P/D ratio is treated as equal to the pre-singularity ratio. This is exact when $\Delta\theta \to 0$ and becomes less accurate as $\Delta\theta$ grows, since each singularity shifts θ and therefore shifts the dividend growth factors. The closed-form is useful for building intuition about the hedging channel, particularly through the comparison of Γ^{AI} and Γ^{N}. Table 1 reports numerically

exact P/D ratios, computed by iterating the Euler equation over the chain of post-singularity states (see Appendix A). The full derivation of the closed form is also in Appendix A.

The key economic content of Proposition 1 is in the comparison of Γ^{AI} and Γ^{N}. Since $\Delta\theta > 0$, AI stocks' dividends grow faster than aggregate consumption upon a singularity ($\Gamma^{AI} > 1+\eta$), while non-AI stocks' dividends grow more slowly ($\Gamma^{N} < 1+\eta$), shrinking as a share of the economy. Combined with the household's high marginal utility in singularity states (due to displacement, $\phi(1+\eta) < 1$ when ϕ is sufficiently small), AI stocks' payoffs are especially valuable, pushing their P/D ratio above that of non-AI stocks. This is the hedging channel: AI stocks pay off precisely when the household's consumption falls, making them a partial hedge against displacement.

It is immediate from the P/D expressions that the valuation spread widens with displacement severity (decreasing ϕ, which raises marginal utility in singularity states via $\phi^{-\gamma}$) and with singularity probability p for sufficiently risk-averse households (higher p puts more weight on the states where AI stocks' dividend advantage operates). The interaction with extinction risk is less obvious:

Proposition 2 (Extinction attenuation)**.** *The valuation spread is decreasing in extinction probability ξ. For the parameterizations considered, the* ratio $(P^{AI}/D^{AI})/(P^{N}/D^{N})$ *also decreases in ξ, since higher extinction probability reduces the weight on non-extinction states where the valuation divergence occurs.*

Proof. Higher ξ reduces both P/D ratios by increasing the probability of total loss. Writing each P/D ratio as $A^j/(1-A^j)$, where A^j is defined in equation (6), note that A^j can be decomposed as $A^j = B[(1-p)+p(1-\xi)S\,\Gamma^j]$, where $B = \beta(1+g)^{1-\gamma}$ and $S = (1+\eta)^{-\gamma}\phi^{-\gamma}$. Higher ξ scales down the singularity component $p(1-\xi)S\,\Gamma^j$ by the same multiplicative factor $(1-\xi)$ in both assets, leaving the non-singularity component $(1-p)$ unchanged. Since $\Gamma^{AI} > \Gamma^{N}$, the singularity component is larger for AI stocks, so this common multiplicative scaling produces a larger absolute reduction in A^{AI} than in A^{N}. The function $f(A) = A/(1-A)$ is increasing and convex on $(0,1)$, with $f''(A) = 2/(1-A)^3 > 0$. Moreover, the semi-elasticity $f'(A)/f(A) = 1/[A(1-A)]$ is increasing in A for $A > 1/2$, so equal absolute reductions produce larger proportional reductions in $f(A)$ at higher A values. Since $A^{AI} > A^{N}$ (because $\Gamma^{AI} > \Gamma^{N}$), the larger absolute reduction in A^{AI} also occurs at a point of higher semi-elasticity, amplifying the proportional decline. The condition $A^j > 1/2$—equivalently, $P^j/D^j > 1$—is satisfied across all parameterizations in Table 1. Therefore the ratio $[A^{AI}/(1-A^{AI})]/[A^{N}/(1-A^{N})]$ decreases in ξ. □

2.3 Discussion

The model's mechanism parallels Gârleanu, Kogan, and Panageas (2012), who show that growth stocks earn lower expected returns because they hedge displacement risk from innovation. In their framework, displacement is driven by new cohorts of firms entering the economy; in ours, it is driven by a discrete AI singularity. The key difference is the nature of the displacement event: GKP model continuous displacement from expanding technological variety, while we model a sudden, severe shift in which the household's consumption falls even as aggregate output rises. This makes the interaction with extinction risk (C. I. Jones, 2024) natural and generates sharper quantitative predictions about how singularity probabilities map to valuation ratios.

The market incompleteness in our model—the household's inability to trade restricted AI equity—is central. If the household could trade the restricted equity, its effective displacement parameter would become $\phi_{\text{eff}} \to 1$: displacement is fully hedged, the SDF no longer overweights singularity states, and the displacement-driven valuation premium largely collapses—though a small residual spread from differential dividend growth ($\Gamma^{AI} \neq \Gamma^{N}$) remains. This echoes Gârleanu, Kogan, and Panageas (2012)'s point that future innovators' rents cannot be traded because the innovators have not yet entered the economy. Our AI owners play an analogous role, though we do not model the entry dynamics that are central to their framework; the displacement in our model comes from the singularity's reallocation of consumption shares rather than from creative destruction by new entrants.

One feature of the discrete singularity that has no analog in Gârleanu, Kogan, and Panageas (2012)'s continuous-displacement framework is that sufficiently severe displacement can violate the existence condition for finite prices (Remark 1). When displacement is extreme, the household's marginal utility in singularity states becomes so large that the SDF-weighted expected dividend growth exceeds the discount rate, and no finite price can clear the market. This discontinuity—from finite to infinite hedging demand—cannot arise under GKP's gradual displacement, where the pricing kernel remains well-behaved. It reflects the distinctive severity of a discrete AI singularity and, as we show in the extensions, motivates the role of government transfers.

3 Quantitative Analysis

Table 1 reports price-dividend ratios across a grid of singularity probabilities and extinction risks. The parameterization uses $\beta = 0.96$, $g = 0.02$, $\gamma = 4$, $\phi = 0.5$ (the household retains half of its consumption share upon displacement, so that $\phi(1 + \eta) = 0.75$ and household

Table 1: Price-Dividend Ratios: AI Stocks vs. Non-AI Stocks

		Price-Dividend Ratio		
p	ξ	AI Stocks	Non-AI Stocks	Ratio
0.1%	0%	10.4	9.8	1.1
0.1%	5%	10.4	9.8	1.1
0.1%	10%	10.3	9.7	1.1
0.1%	20%	10.2	9.7	1.1
0.2%	0%	11.5	10.1	1.1
0.2%	5%	11.4	10.0	1.1
0.2%	10%	11.3	10.0	1.1
0.2%	20%	11.0	9.9	1.1
0.5%	0%	15.5	11.1	1.4
0.5%	5%	15.0	11.0	1.4
0.5%	10%	14.6	10.8	1.3
0.5%	20%	13.8	10.6	1.3
0.8%	0%	21.2	12.3	1.7
0.8%	5%	20.2	12.1	1.7
0.8%	10%	19.2	11.8	1.6
0.8%	20%	17.4	11.4	1.5
1.0%	0%	26.5	13.3	2.0
1.0%	5%	24.8	12.9	1.9
1.0%	10%	23.2	12.6	1.8
1.0%	20%	20.5	12.0	1.7

Parameters: $\beta = 0.96$, $g = 0.02$, $\gamma = 4$, $\phi = 0.5$, $\eta = 0.5$, $\theta = 0.15$, $\Delta\theta = 0.2$. p is the annual singularity probability; ξ is the extinction probability conditional on singularity. AI P/D ratios are numerically exact, computed by iterating the Euler equation over post-singularity states (see Appendix A). The ratio column reports $\mathrm{P/D}^{AI}/\mathrm{P/D}^{N}$.

consumption falls by 25%), $\eta = 0.5$ (aggregate consumption rises by 50% in a singularity), $\theta = 0.15$ (AI stocks are initially 15% of the economy), and $\Delta\theta = 0.2$ (AI's share jumps by 20% of the non-AI remainder upon singularity).

Several patterns emerge. First, AI stock P/D ratios are substantially higher than non-AI stock P/D ratios across the entire grid. For a singularity probability of 0.5% per year with no extinction risk, AI stocks trade at a P/D of about 15.5, while non-AI stocks trade near 11—a ratio of about 1.4. At $p = 1\%$, the ratio rises to 2. Second, increasing the singularity probability raises the AI stock premium: the household values the hedge more as the risk it hedges against becomes more likely. Third, extinction risk compresses the AI premium, as Proposition 2 predicts. At high extinction probabilities, even AI stocks lose value because the states in which they pay off become less likely, reducing both valuations and narrowing

the spread.

The magnitudes are broadly suggestive. As Figure 1 shows, the S&P 500 P/D ratio has reached historically elevated levels, and the AI-heavy NASDAQ has appreciated roughly 50% more than the S&P 500 since 2015, consistent with rising relative valuations for AI-exposed stocks. The mapping from NASDAQ vs. S&P 500 to the model's AI vs. non-AI distinction is imperfect: the NASDAQ is broader than "AI stocks," return differences partly reflect earnings growth rather than valuation multiples, and the S&P 500 itself has substantial AI exposure. Nonetheless, the pattern is consistent with a sustained valuation premium of the kind the model predicts, where AI stock P/D ratios are 1.3 to 2 times those of non-AI stocks across annual singularity probabilities in the 0.5–1% range.

4 Extensions: Market Incompleteness and the Singularity

The baseline model takes market incompleteness as given and studies its pricing implications. This section examines two further consequences: how incompleteness distorts the development of AI, and how government policy might address it.

4.1 Extension 1: Veto and efficient development

We augment the baseline model to allow a positive singularity. With probability p, a singularity occurs. Conditional on a non-extinction singularity (probability $1-\xi$), the singularity is either *positive* (probability q)—the household's share increases to $\alpha_{t+1} = \min(1, \alpha_t/\phi)$ and aggregate consumption jumps by $1+\eta$—or *negative* (probability $1-q$, as in the baseline), with the household's share falling to $\alpha_{t+1} = \phi\,\alpha_t$. We assume $q > 1/2$: the positive singularity is the more likely outcome.

We say AI development is *socially efficient in the Kaldor-Hicks sense*: the total surplus from a non-extinction singularity—summing over both the household and AI owners—is positive, so that the winners could in principle compensate the losers. This holds when $(1+\eta) > 1$, since aggregate consumption rises by a factor $1+\eta$ in both singularity outcomes, and the only question is how the gains are distributed.

The household can *veto* AI development at a cost $\kappa > 0$, representing a permanent fraction of consumption lost to the deadweight costs of intense government intervention needed to halt AI progress. If the household vetoes, it consumes $(1-\kappa)\alpha C_t$ each period with no singularity. We normalize extinction utility to zero; since CRRA utility is negative for all $c > 0$ when $\gamma > 1$, this makes the veto result conservative.

Under complete markets, the household can trade the restricted AI equity—founder stakes, pre-IPO shares, and other claims currently unavailable—so that it fully hedges displacement risk. Its consumption in both singularity states then equals its pre-singularity share of aggregate consumption: $\alpha(1+\eta)C_t(1+g)$.

Proposition 3 (Veto under incomplete markets). *(i) Under incomplete markets: there exists a threshold $\bar{\gamma}$ such that for all $\gamma > \bar{\gamma}$, the household vetoes AI development ($V_{veto} > V_{develop}$), even when development is socially efficient and the veto cost κ is positive.*

(ii) Under complete markets: for κ sufficiently small (in particular, for the same κ as in part (i)), the household never vetoes socially efficient AI development ($V^{CM}_{develop} > V_{veto}$).

Proof. (i) The household's expected one-period utility gain from a non-extinction singularity under incomplete markets is

$$\Delta u(\gamma) = q\, u(\alpha^{+}(1+\eta)) + (1-q)\, u(\phi\alpha(1+\eta)) - u(\alpha), \tag{7}$$

where $\alpha^{+} = \min(1, \alpha/\phi)$ and $u(c) = c^{1-\gamma}/(1-\gamma)$. As $\gamma \to \infty$, the negative-singularity term dominates because $\phi\alpha(1+\eta) < \alpha$ when $\phi(1+\eta) < 1$: the utility cost of the consumption drop grows without bound relative to the utility gain from the positive singularity. Formally, $\lim_{\gamma\to\infty} \Delta u(\gamma)/u(\alpha) = +\infty$ when $\phi(1+\eta) < 1$. In the infinite-horizon problem, $V_{\text{develop}} - V_{\text{veto}}$ is a discounted sum of per-period utility differences; the singularity term contributes $p(1-\xi)\Delta u(\gamma)$ in expectation each period, and since $\Delta u(\gamma) \to -\infty$ while the veto cost κ is fixed, the infinite-horizon difference diverges to $-\infty$ for γ large enough.

(ii) Under complete markets, the household's consumption in both singularity states is $\alpha(1+\eta)C_t(1+g)$. Its expected utility gain is $u(\alpha(1+\eta)) - u(\alpha)$, which is positive since $\eta > 0$. The household strictly prefers development for any $\gamma > 1$ and any κ small enough, and in particular for the same κ used in part (i). □

The veto result illustrates how market incompleteness can distort real decisions, not just asset prices. A numerical example sharpens the point. Using the baseline displacement parameters ($\phi = 0.5$, $\eta = 0.5$, $\xi = 5\%$) with $p = 1\%$, $\gamma = 10$, $\alpha = 0.70$, $q = 0.70$, and a veto cost $\kappa = 1\%$ of permanent consumption: the household vetoes under incomplete markets ($V_{\text{veto}} > V_{\text{develop}}$), even though the positive singularity is more than twice as likely as the negative one. The computation solves the infinite-horizon Bellman equation, treating the singularity as a one-shot event: each period the singularity occurs with probability p, and after realization the economy continues with deterministic growth at the new consumption share. Under complete markets with the same parameters, the household does not veto

($V^{CM}_{\text{develop}} > V_{\text{veto}}$), since it can trade claims on AI capital and consume $\alpha(1+\eta)C_t(1+g)$ in both singularity states. Allowing repeated singularities would reinforce the veto motive by compounding displacement risk.

Extinction risk interacts with this distortion: under our conservative normalization ($U_{\text{ext}} = 0$), higher ξ reduces the weight on non-extinction singularity states, which are the states driving the veto. With a more severe extinction penalty (as C. I. Jones (2024) argues is appropriate for bounded utility), higher ξ would amplify the veto incentive. This connects to debates about AI regulation: Proposition 3 implies that calls to slow or halt AI development may partly reflect unhedgeable downside risk from displacement—the veto is a rational response to the consumption loss that markets cannot insure, not merely technophobia.

C. I. Jones (2024) observes that attitudes toward AI risk depend on consumption levels: wealthier, more risk-averse agents value their current living standards highly and are less willing to accept existential gambles, while poorer agents with higher marginal utility of consumption are more willing to take such risks. Our model offers a complementary channel through financial markets rather than existential risk. Under incomplete markets, households with larger consumption shares—more exposed to displacement—have stronger veto incentives, because the unhedgeable downside looms larger relative to the upside they cannot fully capture. This connects wealth heterogeneity to attitudes about AI development through the displacement channel, not just through differential valuations of life.

4.2 Extension 2: Government transfers

Extension 1 showed that market incompleteness can cause the household to veto socially efficient AI development. A natural question is whether policy can restore efficiency. The ideal solution—broader trading of AI capital—faces the same constraint Gârleanu, Kogan, and Panageas (2012) identify: much of the displacing capital does not yet exist. Government transfers offer an alternative, and the singularity setting gives them a distinctive advantage: if the productivity jump is large enough, the abundance of post-singularity resources can overwhelm the deadweight costs that normally make transfers ineffective. Transfers thus serve a dual role: they reduce AI stock valuations by cushioning displacement (a pricing effect) and, by attenuating the household's unhedgeable downside, they can eliminate the veto distortion (a real effect).

Gârleanu, Kogan, and Panageas (2012) note, in the context of a robustness argument, that intergenerational transfers—bequests, government mandates—would affect the magnitude but not the functional form of the displacement factor; in the limiting case of an altruistic dynasty, bequests eliminate displacement entirely. We take this observation to

the specific setting of an AI singularity, where the key question is whether explosive output growth can overcome the deadweight costs that normally make transfers ineffective. Because the displacing equity may not yet exist—it belongs to future cohorts of innovators—direct trading is infeasible, and government transfers offer an alternative. But transfers ordinarily incur deadweight costs (waste, fraud, administrative burden) that scale with transfer size, making them unattractive.

We model this as follows. The government imposes a tax rate $\tau \in [0, 1)$ on AI owners' post-singularity consumption and transfers the proceeds to the household. Deadweight costs consume a fraction $\delta\tau$ of the transferred amount, where $\delta > 0$ governs the severity of waste. The household's post-transfer consumption in a non-extinction singularity state is:

$$c^H_{post} = \phi\,\alpha\,(1+\eta)\,C_t\,(1+g) + \tau\,(1-\delta\tau)\,(1-\phi\alpha)\,(1+\eta)\,C_t\,(1+g). \tag{8}$$

The first term is the household's displaced consumption. The second is the net transfer: a fraction τ of the AI surplus, reduced by the deadweight cost $\delta\tau$. Note that $(1-\phi\alpha)$ represents the AI owners' share of post-singularity aggregate consumption; this expression uses the household's consumption share α rather than the AI dividend share θ, because the transfer is levied on the consumption allocation, not on publicly traded dividends.

To compute AI stock valuations under transfers, factor equation (8) as $c^H_{post} = \phi_{\text{eff}} \cdot \alpha \cdot (1+\eta)(1+g)C_t$, where

$$\phi_{\text{eff}} \;=\; \phi + \frac{\tau(1-\delta\tau)(1-\phi\alpha)}{\alpha} \tag{9}$$

is the effective displacement parameter. Equation (9) follows directly from dividing (8) by $\alpha(1+\eta)(1+g)C_t$: the first term contributes ϕ and the transfer term contributes the remainder. Since ϕ_{eff} enters the SDF in the same way as ϕ, the P/D formula from Proposition 1 applies with ϕ replaced by ϕ_{eff}. For simplicity, we compute ϕ_{eff} at the initial share α_0; after successive singularities, α declines, slightly altering the effective transfer. This approximation does not affect the qualitative conclusions.

In standard settings (moderate η), the deadweight costs eat into the transfer and the household gains little. But in a singularity with large η, aggregate output grows enormously. The transfer base—AI owners' surplus—grows with η, while the deadweight cost rate is fixed. The ratio of post-transfer to pre-transfer household consumption in the singularity state is independent of the productivity jump η:

$$\frac{c^H_{post}}{c^H_{no\text{-}transfer}} = 1 + \frac{\tau(1-\delta\tau)(1-\phi\alpha)}{\phi\alpha}. \tag{10}$$

This ratio exceeds one whenever $\tau > 0$ and $\delta\tau < 1$ (i.e., deadweight costs do not consume the

entire transfer): transfers always improve the household's position in the singularity state, regardless of η. The economic content is in the *levels*: as η grows, both c^H_{post} and $c^H_{no\text{-}transfer}$ grow without bound, so even inefficient transfers deliver arbitrarily large consumption gains relative to the pre-singularity baseline. To illustrate: under the large singularity ($\eta = 9$, $\phi = 0.05$), raising the deadweight cost parameter to $\delta = 0.9$—so that net transfers per dollar taxed are only $\tau(1-\delta\tau)$, e.g., 0.219 at $\tau = 0.30$—still yields a consumption multiple of 3.5× at $\tau = 0.30$, compared to the 0.5× catastrophe without transfers. Even grossly inefficient redistribution transforms a 50% consumption loss into a 250% gain when the resource base is large enough.

Figure 2 illustrates with two panels, using the baseline risk aversion $\gamma = 4$, $\alpha = 0.70$, $p = 0.5\%$, and $\xi = 5\%$. The baseline uses $\eta = 0.5$ and $\phi = 0.5$; the large singularity uses $\eta = 9$ with $\phi = 0.05$, reflecting the intuition that a singularity powerful enough to produce ten-fold growth would also displace far more labor. The right panel shows that absent transfers ($\tau = 0$), the household faces a catastrophe: consumption halves under the large singularity ($\phi(1+\eta) = 0.5$) and falls by 25% under the baseline ($\phi(1+\eta) = 0.75$). As τ increases, the large singularity responds dramatically—the enormous output growth swamps the deadweight costs—while the baseline gains only modestly.

The left panel shows the corresponding effect on AI stock valuations: transfers reduce the hedge value of AI stocks, compressing P/D ratios. Under the large-singularity parameters, the P/D ratio is undefined at $\tau = 0$: the existence condition in Remark 1 is violated because the household's marginal utility in the singularity state ($\phi^{-\gamma} = 160{,}000$) is so extreme that the pricing sum diverges. In economic terms, the hedge value of AI stocks becomes infinite—no finite price can compensate for the catastrophic displacement. As τ increases and transfers cushion the displacement, effective displacement falls, the existence condition is restored, and finite P/D ratios emerge. The transition from infinite to finite prices illustrates the severity of the incomplete-markets problem and the potential value of transfers.

Figure 2: Government Transfers and the Singularity. Panel (a) shows how transfers compress AI stock P/D ratios by reducing the household's hedging demand. Under the large-singularity parameterization, the P/D ratio is undefined at low tax rates because the existence condition (Remark 1) is violated: displacement is so severe that the household's marginal utility makes the pricing sum diverge, and no finite price can clear the market. Transfers restore finite prices by cushioning displacement. Panel (b) shows the household's consumption change in the singularity state: absent transfers, the household faces a catastrophe (marked at $\tau = 0$), but under the large singularity, even modest tax rates produce enormous consumption gains as explosive output growth overwhelms deadweight costs. Parameters: $\alpha = 0.70$, $p = 0.5\%$, $\xi = 5\%$, $\delta = 0.5$.

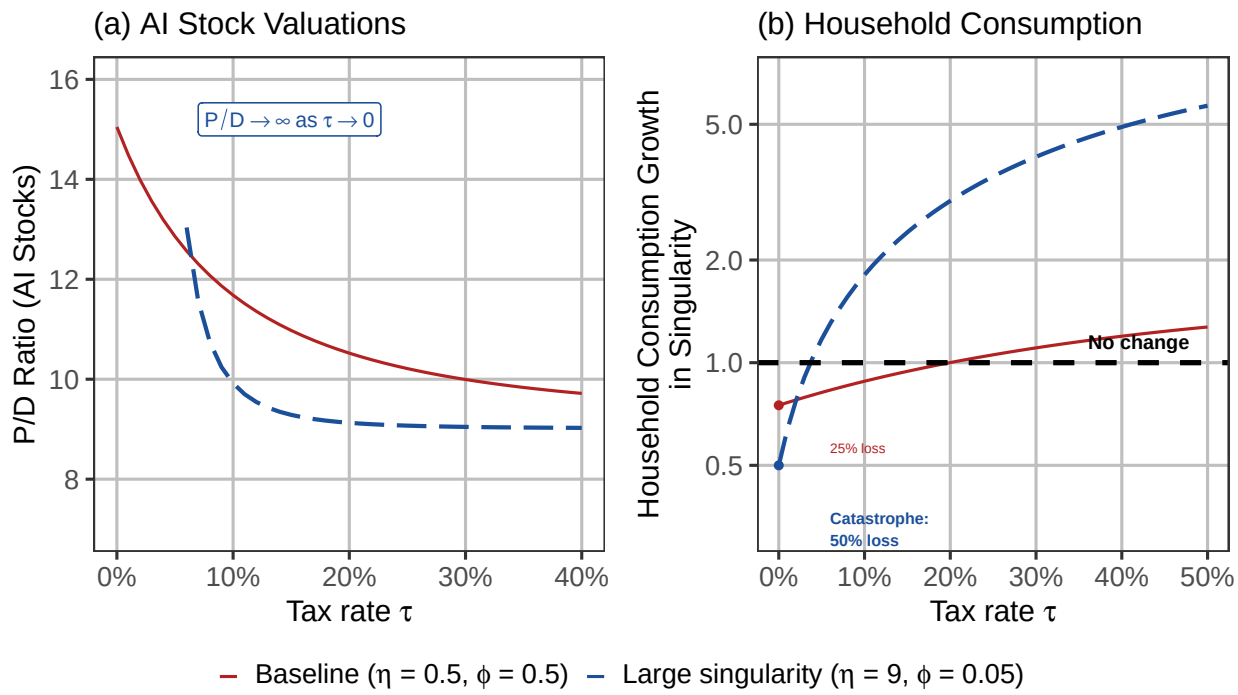

The policy implication is nuanced. In normal times, government transfers are a blunt and wasteful tool for addressing displacement risk. But if an AI singularity produces the kind of explosive output growth modeled by C. I. Jones (2024) and examined critically by Nordhaus (2021), the calculus changes: the abundance of resources makes even inefficient redistribution effective. This suggests that contingent transfer policies—triggered by a singularity—may be worth designing in advance.

5 Conclusion

AI stocks trade at high valuations. We have argued that part of this premium reflects a hedging motive: investors use AI stocks to partially insure against an AI singularity that

would displace their consumption. This mechanism requires market incompleteness—the inability to trade restricted AI equity—and is attenuated by extinction risk, which reduces the weight on the non-extinction states where the hedging channel operates. When markets are incomplete, risk-averse households may inefficiently block AI development, and government transfers, normally wasteful, can become effective if singularity-driven growth is large enough to overwhelm deadweight costs.

Our model is deliberately simple. It abstracts from continuous-time dynamics, heterogeneous beliefs, production-side details, and many other features that would enrich the analysis. The goal is not to provide a definitive account of AI stock valuations but to highlight a specific channel—hedging against displacement under incomplete markets—that connects asset pricing theory to the distinctive features of the AI singularity.

A Proof of Proposition 1

The household's Euler equation for any asset j with price P_t^j and dividend D_t^j is:

$$P_t^j = \mathbb{E}_t \left[\beta \left(\frac{c_{t+1}^H}{c_t^H} \right)^{-\gamma} (P_{t+1}^j + D_{t+1}^j) \right]. \tag{11}$$

Consider the AI stock. Write the price-dividend ratio as $v^{AI} = P^{AI}/D^{AI}$, which is constant in the stationary equilibrium before any singularity has occurred. The household's consumption growth c_{t+1}^H/c_t^H takes three values:

- No singularity (prob $1-p$): $c_{t+1}^H/c_t^H = 1+g$, and $D_{t+1}^{AI}/D_t^{AI} = 1+g$.
- Non-extinction singularity (prob $p(1-\xi)$): $c_{t+1}^H/c_t^H = \phi(1+g)(1+\eta)$, and $D_{t+1}^{AI}/D_t^{AI} = \Gamma^{AI}(1+g)$.
- Extinction (prob $p\xi$): $c_{t+1}^H = 0$, so the payoff is zero.

Substituting into the Euler equation with $P_t^{AI} = v^{AI} D_t^{AI}$:

$$\begin{aligned} v^{AI} D_t^{AI} = \beta(1+g)^{-\gamma} \Big[& (1-p)(1+g)(v^{AI}+1) D_t^{AI} \\ & + p(1-\xi)[\phi(1+\eta)]^{-\gamma}(1+g)\, \Gamma^{AI}(v^{AI}+1) D_t^{AI} \Big]. \end{aligned} \tag{12}$$

Note that after a non-extinction singularity the economy resets with new parameters (α', θ'), and the *new* P/D ratio in the post-singularity state will generally differ because Γ^{AI} depends on θ. (The non-AI growth factor $\Gamma^N = (1-\Delta\theta)(1+\eta)$ is θ-independent, so the

non-AI closed form is exact.) For tractability, we approximate the post-singularity AI P/D ratio by v^{AI}. This approximation becomes exact as $\Delta\theta \to 0$. For the numerically exact values reported in Table 1, we solve the Euler equation by backward recursion over the chain $\theta_0, \theta_1 = \theta_0 + \Delta\theta(1-\theta_0), \theta_2, \ldots$, starting from a terminal state where θ is close to one and the approximation error vanishes. Dividing by D_t^{AI} and solving for v^{AI}:

$$v^{AI} = \frac{\beta(1+g)^{1-\gamma}\left[(1-p) + p(1-\xi)(1+\eta)^{-\gamma}\phi^{-\gamma}\Gamma^{AI}\right]}{1 - \beta(1+g)^{1-\gamma}\left[(1-p) + p(1-\xi)(1+\eta)^{-\gamma}\phi^{-\gamma}\Gamma^{AI}\right]}. \tag{13}$$

This can be rewritten as equation (4). The derivation for non-AI stocks is identical, replacing Γ^{AI} with Γ^{N}.

References

Acemoglu, Daron (2025). "The Simple Macroeconomics of AI". In: *Economic Policy* 40.121, pp. 13–58.

Aghion, Philippe, Benjamin F. Jones, and Charles I. Jones (2019). "Artificial Intelligence and Economic Growth". In: *The Economics of Artificial Intelligence: An Agenda*. University of Chicago Press, pp. 237–290.

Barro, Robert J. (2006). "Rare Disasters and Asset Markets in the Twentieth Century". In: *Quarterly Journal of Economics* 121.3, pp. 823–866.

Gârleanu, Nicolae, Leonid Kogan, and Stavros Panageas (2012). "Displacement Risk and Asset Returns". In: *Journal of Financial Economics* 105.3, pp. 491–510.

Jones, Charles I. (2024). "The AI Dilemma: Growth versus Existential Risk". In: *American Economic Review: Insights* 6.4, pp. 575–590.

Knesl, Jiří (2023). "Automation and the Displacement of Labor by Capital: Asset Pricing Theory and Empirical Evidence". In: *Journal of Financial Economics* 147.2, pp. 271–296.

Kogan, Leonid and Dimitris Papanikolaou (2014). "Growth Opportunities, Technology Shocks, and Asset Prices". In: *Journal of Finance* 69.2, pp. 675–718.

Kogan, Leonid, Dimitris Papanikolaou, and Noah Stoffman (2020). "Left Behind: Creative Destruction, Inequality, and the Stock Market". In: *Journal of Political Economy* 128.3, pp. 855–906.

Nordhaus, William D. (2021). "Are We Approaching an Economic Singularity? Information Technology and the Future of Economic Growth". In: *American Economic Journal: Macroeconomics* 13.1, pp. 299–332.

Pástor, Ľuboš and Pietro Veronesi (2009). "Technological Revolutions and Stock Prices". In: *American Economic Review* 99.4, pp. 1451–1483.

Wachter, Jessica A. (2013). "Can Time-Varying Risk of Rare Disasters Explain Aggregate Stock Market Volatility?" In: *Journal of Finance* 68.3, pp. 987–1035.